Data Integration on High-Difficulty Binary Classification

Julia Finch, Jesse Hellemn, Kentaro Hoffman, Zhe Zhang

May 1, 2017

Abstract

This paper explores the effectiveness of the data integration technique Generalized Multiple Kernel Learning (GMKL) on high-difficuly binary classification data. GMKL that integrates learned kernels from two disparate data sources is systematically compared to performing GMKL on the two sources naively concatenated as well as standard classification algorithms that are performed on the concatenated sources. The variables that are systematically varied are the number of informative dimensions in each source, the relative information provided by each source, and the number of useless noise dimensions added to each source. It is found that GMKL consistently outperforms its competitors. This paper provides evidence that data integration techniques, specifically GMKL, have the ability to drastically improve upon the performance of naive concatenation.

Introduction

For classification tasks in many fields, especially in medicine and the social sciences, it is increasingly common for data to come from multiple disparate data sources. For example, a sociologist might be interested in predicting future income using two different sources, one on family environment and one on school environment (we will use "source" to refer to a single coherent dataset). For the best analysis and classification results, all data sources should be taken into account. However, most current machine learning classification techniques have been developed for only single dataset inputs, and it is not obvious how to best adapt these techniques to multi-source data.

A naive approach is to concatenate the sources together into one large feature matrix, essentially treating all of the data as a single incoherent source. Although simple, this method throws away information about the separate sources and forces data analysis and classification techniques to treat all of the sources in the same way. Various data integration techniques have been proposed to more cleverly and effectively combine multiple sources. Unfortunately, these techniques have been poorly tested, and there has been no systematic evaluation of their effectiveness.

This paper evaluates one such specialized data integration technique, Generalized Multiple Kernel Learning (GMKL), against traditional classifiers that use concatenated data. We simulate many 2-source datasets with a variety of properties for these comparative tests. We will use the term "classifier" to refer to both specialized data integration techniques such as GMKL as well as naive methods that first concatenate all sources together.

Data Generation

This paper follows the systematic philosophy of simulation studies advocated in [NBM14].

Criteria of Good Simulated Data

In order to understand the behavior of the classifiers on 2-source data as properties of the data are varied, we systematically created data to satisfy all of the criteria detailed below. These criteria allow us to test all the classifiers fairly against each other on consistent benchmarks.

Criteria of Simulated Data
1. The dataset consists of two separate sources.
2. The two classes are separated by a true decision boundary that is known and calculable.
    - or the two classes can be specified to overlap with percentage p, where $p = 0$ leads to no overlap and $p = 50$ leads to complete overlap
3. The decision boundary's complexity (and thus difficulty) can be parametrized and controlled.
4. The reliability of each source can be specified.
5. The noisiness of each source can be specified.
    - extra meaningless dimensions can be added to each source

Explanation for Data Criteria

Suppose that a classifier C only obtains 60% classification accuracy on a dataset D (with datapoints evenly split amongst 2 classes). This could be attributable to either:

- The classifier is not well suited to certain properties of dataset D.

- Eighty percent of both classes overlap with each other. The best possible strategy in this area of overlap is to guess the class with 50% accuracy. An optimal classifier will then guess 40% of the datapoints correctly and also classify the 20% of non-overlapping datapoints perfectly.

Criterion 2 ensures that the latter case does not occur, so that classifier performance is attributable solely to its efficacy on certain types of data.

Criteria 5 is more difficult than it at first seems. In order to create a source with N_{useful} useful dimensions and N_{noisy} meaningless dimensions, we needed to both 1) make N_{noisy} dimensions of pure noise and 2) make N_{useful} dimensions, all of which are always useful. With a random coefficient linear model, it is impossible to verify that all N_{useful} dimensions are actually useful.

Data Generation Models

Random Coefficient Linear Model

A simple, common way to simulate data is to use a linear model with random coefficients. This model generates data by:

1. Specify the number of true latent variables K along with the number of visible variables N

2. Specify two distribution D_j^0 and D_j^1 of each latent variable z_j, $1 \leq j \leq K$, where D_j^0 is the distribution of z_j for negative classes and D_j^1 is the distribution of z_j for positive classes

3. Specify how each visible variable x_j is generated from the latent variables z_i with a formula of the form

$$x_l = \sum_{i=1}^{K} \beta_i z_i + \sum_{i=1}^{K} \sum_{j=i}^{K} \beta_{i,j} z_i z_k + \text{higher-order-interactions}$$

of linear combinations of arbitrary functions of the latent variables

4. Specify every β in the above formula

5. For every datapoint

(a) Pick which class the datapoint belongs to

(b) Sample each z_j from its respective distribution for this class

(c) Generate each visible variable x_j from its formula

This model has significant shortcomings

- It is not known how to systematically make the classification problem more or less difficult

- It is hard to know if the generated data overlaps

- It is hard to pick the βs to ensure that all of the criteria in are satisfied

- There are many distributions and formulas to specify arbitrarily

- It is hard to know how many of the generated dimensions are useful

Feed-forward Network Model

In order to satisfy all of the criteria in , we created a feed-forward conditional network (Figure 1). This network's process is given in detail Algorithm 1. The network model was inspired by Bayesian Networks; the output of each node is sampled conditionally after that node's inputs ahave all been calculated. The model is highly extensible. For this particular network, each sources reliability, difficulty, noisiness, and relative amount of noise to useful information could all be systematically varied independently of each other. separately.
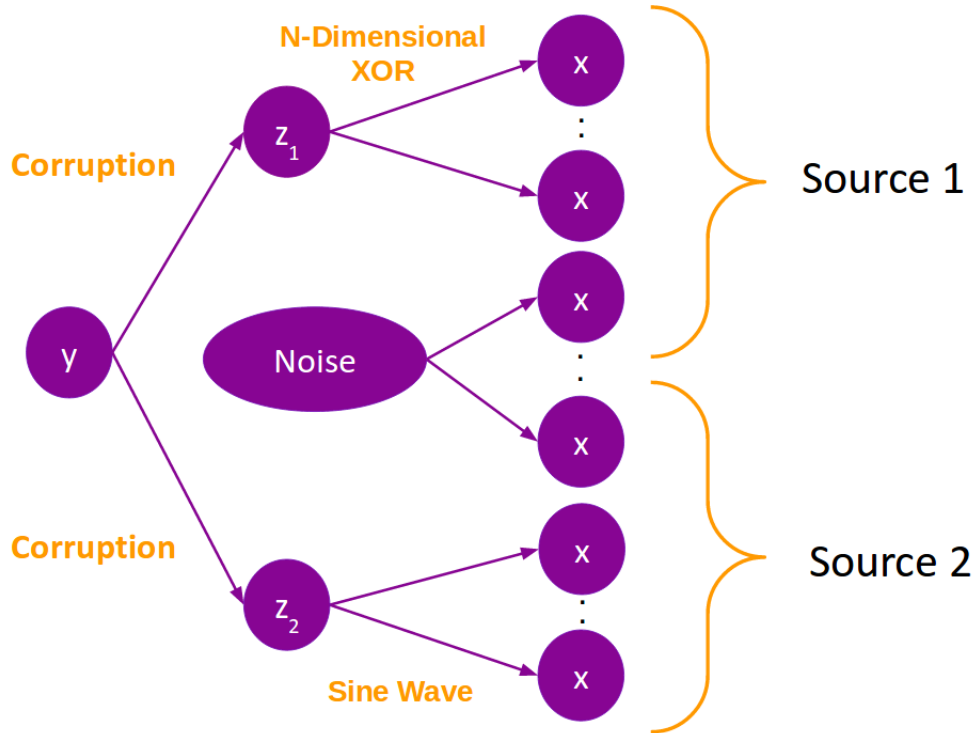

Figure 1: *General schematic of the network data generation model. Every circle is a "node" with takes in an input from the previous layer and outputs a new layer. For example, z_1 will produce a N-dimensional XOR, where the parity depends on the input from y.*

The above model creates two sources of data with very different types of decision boundaries.

The first source is an N-dimensional XOR, which consists of clusters of points at every corner of an N-dimensional binary hypercube, where each corner belongs to a different class than all of its $N-1$ closest

---
Algorithm 1 Data generation process for the network model

    ▷ Sample the y layer
1: $y \leftarrow \text{Bernoulli}(p)$

    ▷ Sample the z layer
2: **for all** $z_i \in \{z_1, z_2\}$ **do**
3:      $c \leftarrow \text{Bernoulli}(p_i)$ ▷ p_i chance to corrupt source i
4:      $z_i \leftarrow c * (1 - y) + (1 - c) * y$ ▷ If corrupting, z_i will be 0 if y is 1 and 1 if y is 0
5: **end for**

    ▷ Sample the x layer of source 1, an N_1-dimensional XOR
6: **if** z_1 is even **then**
7:      $x_1^{(1)}...x_{N_1}^{(1)} \leftarrow N_1$-dimensional binary vector of even parity
8: **else**
9:      $x_1^{(1)}...x_{N_1}^{(1)} \leftarrow N_1$-dimensional binary vector of odd parity
10: **end if**

    ▷ Sample the x layer of source 2, a k_2-period sine wave
11: $x_1^{(2)} \leftarrow \text{Uniform}(-k_2\pi,\ k_2\pi)$
12: $x_2^{(2)} \leftarrow \text{Uniform}(-k_2\pi,\ k_2\pi)$
13: $x_3^{(2)} \leftarrow (x_1^{(2)} + x_2^{(2)})sin(x_1^{(2)}) + mz_2$

    ▷ Add noise
14: $x^{(1)} \leftarrow x^{(1)} + \text{Normal}(0, \sigma_1)$ ▷ Gaussian noise added to each dimension independently for XOR sources
15: $x_3^{(2)} \leftarrow x_3^{(2)} + \text{Normal}(0, \sigma_2)$ ▷ Gaussian noise only added in direction of margin for sine wave sources
---

neighboring corners. When the input to z_1 is a 1 (representing a positive example), then the output is sampled from all binary 0-1 N-dimensional vectors with an even numbers of ones, also known as an even-parity XOR. Figure 2 shows a 3D example. The complexity of the decision boundary is increases with the number of clusters, and the number of clusters increases exponentially with the dimension of the XOR. Noise is added to an XOR by adding sampling from Normal(0, σ_1) independently for every dimension.

The second source is a sine wave in 3D space (Figure 3); the complexity of this decision boundary is directly proportional to the period of the wave and inversely proportional to the margin m between the two classes. Noise is added to the sine wave only in the direction of the margin (points can only be moved "up" or "down").
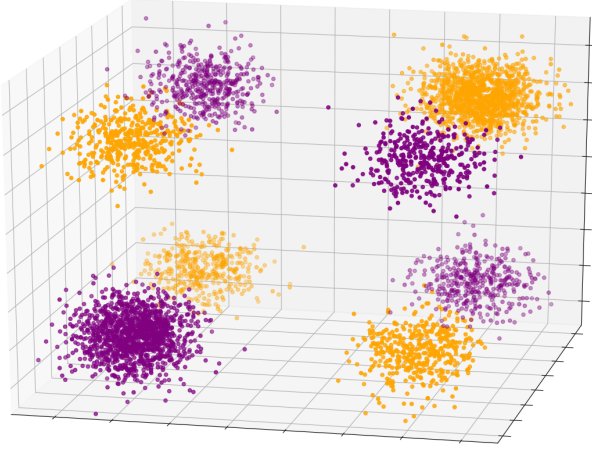

Figure 2: *A 3D-XOR with equal class size and $\sigma_1 = 0.2$ noise. The noise for the XORs created by this model are independent Normal(0, σ_1) added to every dimension.*
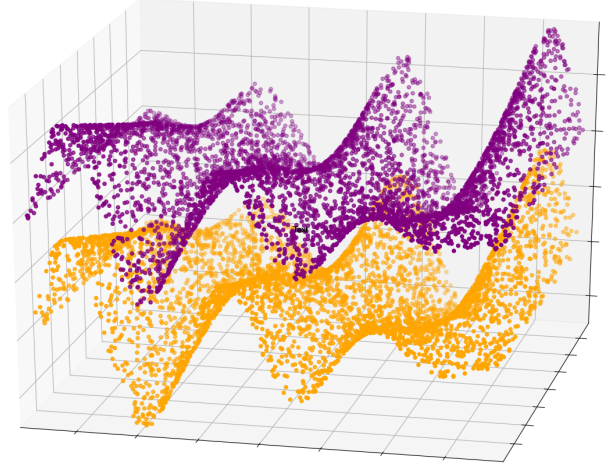

Figure 3: *A 3 period sine wave with $\sigma = 0$ noise and a margin of $m = 10$. The equation of this wave in xyz coordinates is $z = (x + y)sine(x) + mc + Normal(0, \sigma_2)$, where m is the margin and c is the 1 or 0 class label.*

This model has several nice properties:

- For both the N-dimensional XOR and the sine wave (with small enough margin m), every dimension is necessary for perfect classification.

- The complexity of the decision boundary for each source can be easily controlled.

- The reliability of each source can be varied independently of the other.

- The data can be made more noisy (to an extent) without compromising the separability of the two classes.

- Extra meaningless dimensions can easily be added to either source.

The reliability of each source is controlled by corrupting its input signal by a speficied percentage. For example, to decrease the reliability of source 1 to 80%, p_1 is set to 0.1. Then z_1 will only equal y 90% of the time, and will give the wrong signal 10% of the time. Note that 0% reliability is achieved when the signal is corrupted 50% of the time, since the signal is a binary 0 or 1.

The first experiments in this paper use a similar model to the one above, except with another XOR as the second source instead of the sine wave. We refer to this model as a "Double XOR." The Double XOR is not equivalent to one big XOR. For example, a 10-dimensional XOR has $2^{10} = 1024$ clusters while a Double XOR with two 5-dimensional XORs only has $2^5 * \frac{2^5}{2} = 512$ clusters, since a even parity XOR in one source implies an even parity XOR in the other source (if neither source is corrupted). More importantly, it is no longer necessary to have all dimensions to know the parity of the XOR.

The final experiment uses the exact model described above.

Classifiers

The data integration method GMKL was tested against standard machine learning classifiers Gaussian Naive Bayes, Random Forest, and an SVM using a RBF kernel. The machine learning classifiers were fed concatenated data, and so ignored separate sources. Additionally, standard GMKL was compared against GMKL that was fed concatenated data, to isolate the effects of using a separate kernel for each data source. A brief overview of every classifier and of GMKL follow.

Where possible, Python's sklearn implementations were used. Classifiers with hyper parameters, such as the SVM, were trained using cross-validation over a suite of possible parameters, as is industry standard.

Notation

In describing different classification techniques, the following notation will be used:

- (x, y) pair denotes a feature vector x $\in \mathbb{R}^m$ and its corresponding target value $\in \{0, 1\}$.

- x_j denotes the jth feature in x.

- there are N training examples and ith training pair is represented by (x^i, y^i) .

Classifiers Using Concatenated Data

Gaussian Naive Bayes

Model: Gaussian Naive Bayes assumes that every feature (dimension) x_j is generated independently conditional on the value of y i.e. $P(x|y) = \prod_{j=1}^{m} P(x_j|y)$. In addition it also assumes that each $P(x_j|y) \sim \mathbf{N}(\mu_{\mathbf{j},\mathbf{y}}, \sigma_{\mathbf{y}})$. To predict $\hat{y}$ from a given x, the model computes the posterior probability $P(y|x)$ and pick the more likely case. This is not a very powerful model.

Properties: The formulation is clean and simple. Naive Bayes is very fast to train and is somewhat resistant to extra noise dimensions. However, the assumption of independence between dimensions is very strong, and means that only 2nd degree decision boundaries can be fit. It is expected to perform very poorly on the complex, very non-linear Double XOR dataset.

Implementation: Gaussian Naive Bayes does not have any hyper parameters and did not need any parameter tuning.

Random Forest

Model: Random Forest uses an ensemble of deep decision trees. Each decision tree is based by a bootstrap sample of the training data, and the splitting nodes are constrained to a random subset of all features. In this manner, overfitting of individual decision tree is controlled by ensembling.

Properties: Random Forest is known to be very resistant to extra noise dimensions, since uninformative features are simply never used to split decision tree. Random Forests make no assumptions about the input data, and are unaffected by scaling or large differences in scale between features.

Implementation: We used cross validation to select the optimal number of trees $\bar{t} \in \{10, 100\}$. Other parameters were the defaults in the Python sklearn implementation: the gini index was used as the measure quality of the tree, up to $\sqrt{num\_features}$ features were used for each tree, and trees are split until all nodes are pure or have ≤ 2 points.

K-Nearest Neighbors (KNN)

Model: KNN labels a given point x as the majority vote label of its K-closest points (in the training data) i.e.

$$f(x) = sign(\sum_{x^i \in \text{K closest points}} y^i)$$

Properties: KNN is able to fit extremely complex decision boundaries. However, it often suffers from the curse of high dimensionality; the distance between points is distorted in high dimensions. Importantly, KNN is very affected by meaningless noisy features, as these features factor into the Euclidian distance between points just as much as the useful features.

Implementation: We used cross validation to select the best number of neighbors $\bar{k} \in \{1, 2, 10\}$

Support Vector Machine (SVM) with a Radial-Basis-Function (RBF) Kernel

Model: SVMs find the largest-margin between two classes, using a kernel function to first project the data into a higher-dimensional space. Similar to KNN, SVMs label data points considering the weighted voting of a set of similar data points

$$f(x) = sign(\sum_{i=1}^{N} y^i d^i K(x, x^i)) \text{ where K(x, } x^i) = exp(-\gamma |x - x^i|^2)$$

where d^i is learned from data by finding the largest-margin separating hyperplane in the projected space.

Properties: SVMs can fit very complex decision boundaries, and don't suffer from the curse high dimensionality as much as KNN does. The complexity of the decision boundary learned by the SVM can be controlled with two hyper parameters γ and C (kernel width and regularization strength respectively). These characteristics make SVM the go-to algorithm for many complex classification problems. However, just like KNN, it is unable to effectively discriminate against extra noise dimensions. SVMs also require lots of data points to learn complex decision boundaries.

Implementation: Cross-validation was performed to search for the best regularization parameter $\bar{C} \in \{0.1, 1, 10, 100\}$ and kernel width $\bar{\gamma} \in \{.01, .1, 1, 10\}$.

Generalized Multiple Kernel Learning (GMKL)

GMKL learns separate kernels for each data source before integrating the separate kernels in an optimization step that's based on a global error measure; the resulting combined kernel is then fed into an SVM. Although this method uses a SVM for actual classification, the separate kernels allow each source to have its own representation. A brief description of the algorithm is described below, with a general flowchart of the process shown in Figure 4 and a more detailed explanation of the optimization part in Algorithm 2. A full explanation of the GMKL algorithm can be found in [YXY+11].

Model: The final trained model is almost the same as those from SVM, except that the final kernel function is actually learned from data. More specifically, the final kernel is a convex combination of a set of predefined kernel functions:

$$K(x, x^i) = \sum_{q=1}^{Q} \theta_q * K_q(x, x^i) \text{where } \theta_q \text{is learned from data.}$$

7

Formulation as Optimization Problem

$$\underset{\theta,v,b}{\text{minimize}} \qquad C\frac{1}{N}\sum_{i=1}^{N}L(f_{\theta,w,b}(x^i),y^i) + 1/2\sum_{q=1}^{Q}|\frac{v_q|^2}{\theta_q}$$

$$\text{subject to} \qquad \beta|\theta|_2^2 + (1-\beta)|\theta|_1 \le 1$$

Where C is the regularization parameter, β is the elastic net parameter, L is the hinge loss function and

$$f_{\theta,w,b}(x^i) = \sum_{q=1}^{Q} w_q\phi_q(x^i)\sqrt{\theta_q} + b$$

Now if we define v by $v_q := \frac{w_q}{\sqrt{\theta_q}}$, then we have a convex optimization problem:

$$\underset{\theta,v,b}{\text{minimize}} \qquad C\frac{1}{N}\sum_{i=1}^{N}L(f_{\theta,w,b}(x^i),y^i) + 1/2|w|^2$$

$$\text{subject to} \qquad \beta|\theta|_2^2 + (1-\beta)|\theta|_1 \le 1$$

$$\text{Where} \qquad f_{\theta,v,b}(x^i) = \sum_{q=1}^{Q} v_q\phi_q(x^i) + b$$

Observe that the elastic net constraint ensures sparsity and grouping effect for the set of kernels chosen for the final model.

Optimization Algorithm

Algorithm 2 Level method for the MKL[1]

    ▷ Initialization
1: $t \leftarrow 0$
2: Let θ be uniformly initialized subject to the elastic net constraint

3: **while** difference between the upper bound and lower bound is less than ϵ **do**
4:     $\alpha^t = \underset{\alpha}{argmax}D(\theta^{t-1},\alpha)$ ▷ Solve dual problem
5:     $h^t(\theta) = \underset{1\le i\le t}{max}D(\theta,\alpha^i)$ ▷ Construct a cutting plane model
6:     Calculate a lower bound and an upper bound for the optimal solution $\overline{D_t}, \underline{D_t}$ and an improvement
    set level set $L = \{\theta : h^t(\theta) \le$ some convex combination of $\overline{D_t}, \underline{D_t}\}$
7:     Project θ^{t-1} to L to obtain θ^t
8: **end while**

Implementation For each 'independent' data source, we constructed 10 RBF kernels with the width $\gamma \in \{2^{-3}, 2^{-2}, 2^{-1}..., 2^6\}$. Then we trained our model with the regularization parameter C fixed to 100.
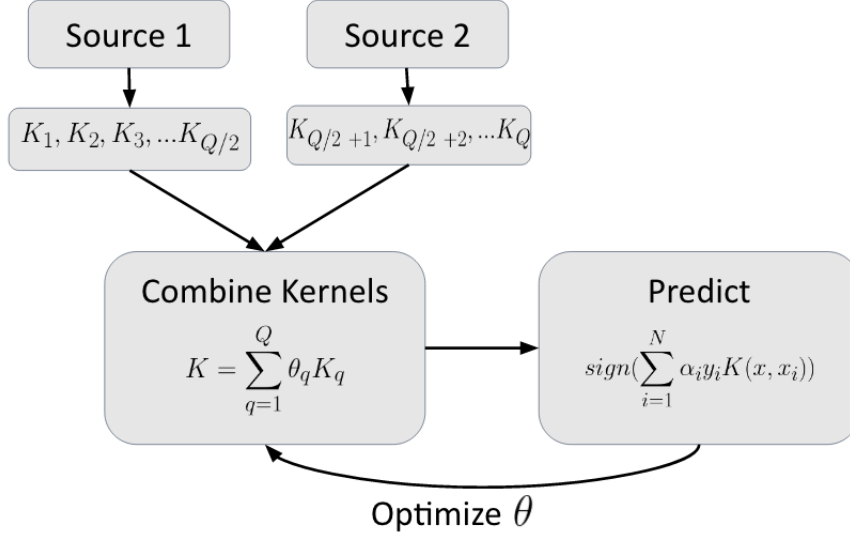
Figure 4: *Flow Chart of Our Model*

Concatenated GMKL To isolate the effect of data integration, we also applied GMKL to concatenated data for comparison. More specifically, we constructed only 10 RBF kernels with $\gamma \in \{2^{-3}, 2^{-2}, 2^{-1}..., 2^6\}$ for the concatenated dataset and then trained GMKL algorithm with C fixed at 100.

Experiments

Several experiments were run. A full description of parameters can be found in the Appendix.

Experiment 1: Data Dimension Scaling

The first experiment used the network data generation model shown in Figure 1, using an XOR for both sources. Experiment parameters are given in Table 1 and results in Figure 5. This first experiment measured classification accuracy of GMKL, Concatenated GMKL, SVM, KNN, and Random Forest as the dimensionality of the XOR (and thus the difficulty of the classification task) increased.

This test increases to a maximum of 7 dimensions per source. At first glance this may seem like a very small problem, but recall that the XOR has a very complex decision boundary, as evidenced by every classifier's poor performance at even this low XOR dimensionality.

$M_{i,useful}$	$M_{i,noisy}$	N	T	p_1, p_2	σ_i	C	k
[2,3,4,5,6,7]	0	5000	1:2	0,0	0.2	1:1	NA

Table 1: *Full data generation parameters for Experiment 1. The two sources always had the same dimension as each other. Noise of variance $\sigma_i = 0.2$ corresponds to nearly touching but still separable clusters.*
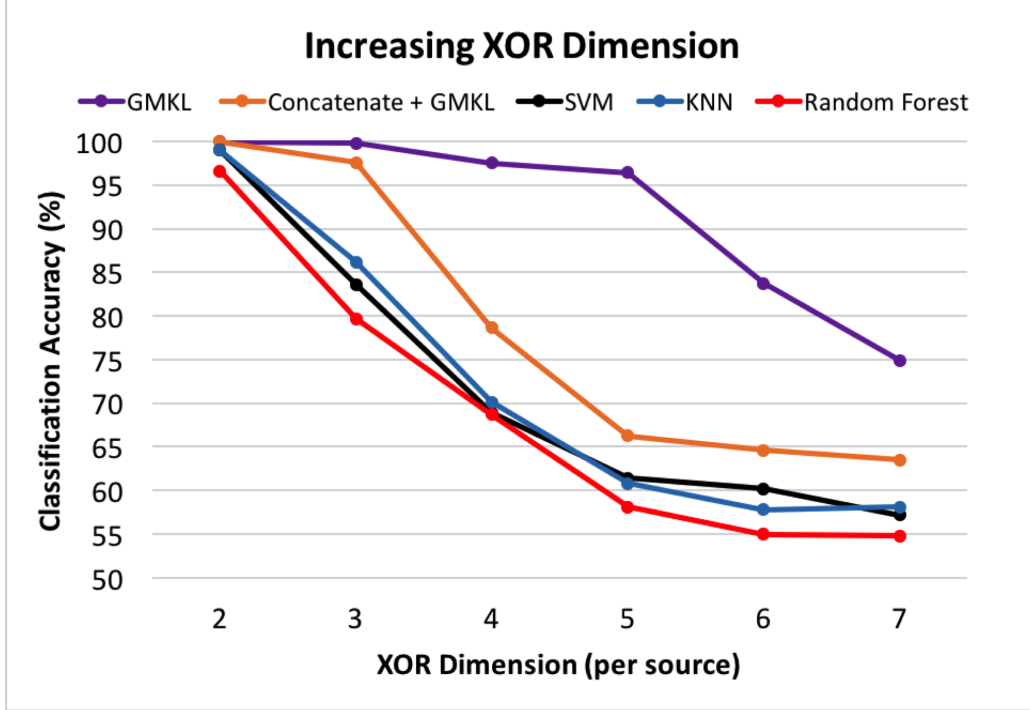
Figure 5: *Classifier accuracies as a function of increasing XOR dimension (increasing problem complexity). GMKL using separate kernels for each source outperforms all other classification methods when dimension increases.*

All classifiers performs worse as the dimension increases. At 7 dimensions each, there are 2^{13} clusters of about $\frac{5000}{2^{13}} < 1$ data points each, which is a very difficult classification problem. GMKL is so much more effective than the other classifiers because it treats each source sepparately, effectively cutting the dimensionality of the problem in half. When both sources are 7-dimensional, GMKL builds kernels on 7-dimensional datasets rather than on the concatenated 14-dimensional dataset. Thus, GMKL on $2n$-dimensional seems to perform comparably to other classifiers on n-dimensional data.

Experiment 2: Corrupted XOR Sources

The second experiment investigated the effect of unreliable sources on classification accuracy. Sources were made unreliable by corrupting their signal in a specified percentage of generated data. Effective data integration techniques should be able to rely on the more reliable source and ignore unreliable information.

The reliability of each source was varied from 100 to 0% in increments of 20% (by varying the corruption probabilites from 0 to 50% in increments of 10%). Classification accuracies for GMKL with separate kernels and GMKL on concatenated data are shown in Figure 6 and 7. This experiment was repeated for XOR dimensions of 3, 5, and 7. Only the results when XOR dimensions was 5 are presented here; the full results can be found in the Appendix. Note that at 50% corruption of each source, the true signal is completely lost in noise.

$M_{i,useful}$	$M_{i,noisy}$	N	T	p_1, p_2	σ_i	C	k
[3,5,7]	0	5000	1:2	[0.0,0.1,0.2,0.3,0.4,0.5],[0.0,0.1,0.2,0.3,0.4,0.5]	0.2	1:1	NA

Table 2: *Full data generation parameters for Experiment 2. The two sources still always had the same dimension as each other. Corruption levels of $p_1 = u, p_2 = 2$ is symmetric to $p_1 = y, p_2 = x$ since both sources are XOR, so each of these combinations was run only once.*

GMKL

Source 2 Corruption (%) \ Source 1 Corruption (%)	0	10	20	30	40	50
0	94.72	92.32	88.84	88.6	91.24	91.06
10	92.32	85.53	80.74	79.18	81.57	79.72
20	88.84	80.74	68.27	71.09	63.11	59.57
30	88.6	79.18	71.09	63.59	62.42	56.63
40	91.24	81.57	63.11	62.42	53.66	52.19
50	91.06	79.72	59.57	56.63	52.19	48.47

Concatenate + GMKL

Source 2 Corruption (%) \ Source 1 Corruption (%)	0	10	20	30	40	50
0	67.67	65.75	60.38	58.61	60.83	55.97
10	65.75	62.45	61.25	56.87	54.95	60.44
20	60.38	61.25	60.95	55.43	56.51	54.89
30	58.61	56.87	55.43	54.26	53.99	52.28
40	60.83	54.95	56.51	53.99	49.97	48.35
50	55.97	60.44	54.89	52.28	48.35	48.83

Figure 6: *GMKL's classification accuracy for all combinations of source corruptions. GMKL's performance only drops when both sources are corrupted, as it is able to rely on the more reliable source.*

Figure 7: *GMKL's classification accuracy for all combinations of source corruptions when fed concatenated data. Performance is noticebly worse, as GMKL can no longer ignore the unreliable source and is much more easily confused.*

GMKL performs much better when it is allowed to create separate kernels for each data source, since then it can easily ignore the more unreliable source. This is a promising result for data integration in general, and not just for GMKL. Treating each source separately allows for only the most useful information from each source to be utilized without confusion from the other source.

Experiment 3: Noise Dimensions

The third experiment investigated the effect of extra, meaningless, noisy dimensions on classification accuracy. Extra dimensions of pure Gaussian noise that provided no information on the true class label were added to each source. Good classifiers should be able to disregard the useless dimensions and retain good accuracy. Results are shown in Figure 8.

This experiment was repeated three times, for 3, 5, and 7 XOR dimensions per source. The percentage of dimensions that are uninformative is varied from none to a quarter, third, half, and then finally two thirds noisy uninformative dimensions. The purpose of measuring the proportion of noise variables as opposed to the number of noise variables is so that the three experiments over different XOR dimensions can be accurately compared.

The results from the three dimensional Double XOR experiment are represented in Figure 8. The results from the five dimensional and seven dimensional Double XOR can be found in the Appendix.

$M_{i,useful}$	$M_{i,noisy}$	N	T	p_1, p_2	σ_i	C	k
3	[0,1,2,3,6]	5000	1:2	0,0	0.2	1:1	NA
5	[0,1,3,5,10]	5000	1:2	0,0	0.2	1:1	NA
7	[0,1,4,7,14]	5000	1:2	0,0	0.2	1:1	NA

Table 3: *Full data generation parameters for Experiment 3. The sources always had the same number of useful dimensions and noisy dimensions as each other.*
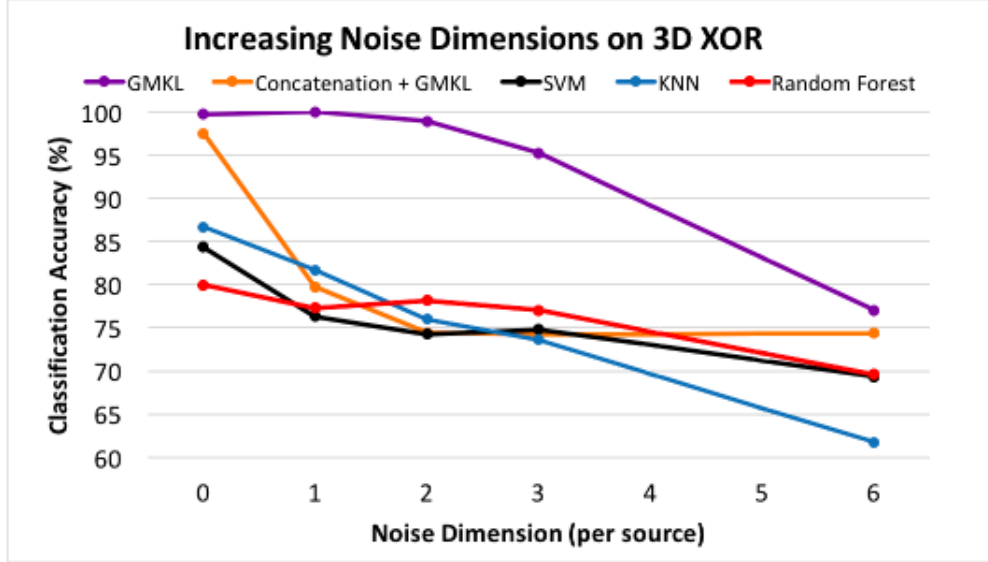
11

Figure 8: *Classification accuracy as a function of increasing number of uninformative dimensions per source. Again, GMKL outperforms all other classifiers by treating each source separately.*

There are a few interesting things to note of these results.

- Random Forest is least affected by the addition of uninformative dimensions.

- KNN is most affected by the addition of uninformative dimensions.

- GMKL is only resistant to added uninformative dimensions when it is allowed to construct separate kernels for each source. When it is forced to use concatenated data, it fares just as well as the other classification techniques.

The only classifier that is able to both get good accuracy and stay resistant to added noise dimensions is GMKL using separate kernels for each source. By treating each source separately, it is able to ignore much of the extra noise introduced by the second source. Instead of building a kernel for a 3-dimensional XOR + another 3 dimensions of a redundant XOR + 6 total dimensions of uninformative noise, it can build a kernel to handle just a 3-dimensional XOR with 3 uninformative dimensions.

Experiment 4: Corrupted Sine and XOR Sources

This experiment did not use the Double XOR, but had a sine wave as the second source exactly as in Figure 1. This experiment is thus much more similar to real multi-source data, as the two sources have different types of decision boundaries and very different properties.

Corruption probabilities were varied as in Experiment 2, but in larger increments. Results for a 3D XOR and 2-period sine wave are presented in Figure 9.

$M_{i,useful}$	$M_{i,noisy}$	N	T	p_1, p_2	σ_i	C	k
3	0	0	1:2	[0.0,0.15,0.3,0.45],[0.0,0.15,0.3,0.45]	0.2	1:1	[1,2,3]

Table 4: *Full data generation parameters for Experiment 4. Corruption probabilites were varied in increments of 15% instead of 10%, but the trends are just as visible. The period of the sine wave affected classification accuracies but did not affect any trends as the corruption probabilites changed.*
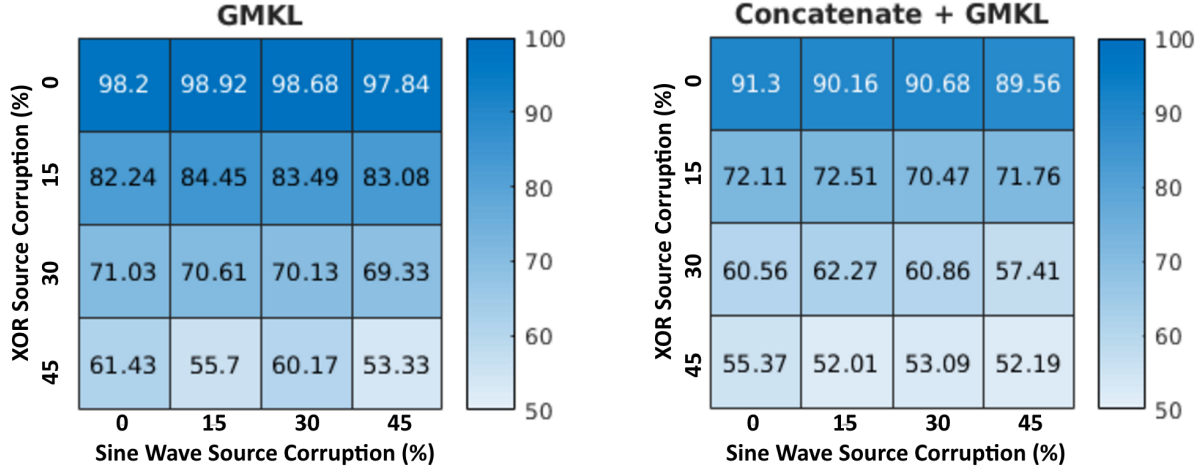
Figure 9: *Classification accuracies for GMKL (with both separate kernels and with concatenated data) for different corruption probabilites. For both cases, GMKL relies on the XOR and ignores the sine wave, though the GMKL with separate kernels is more able to ignore the sine wave source.*

As can be seen by GMKL's insensitivity to corruption of the sine wave source, GMKl is relying on the XOR source and does not handle the sine wave data well (this was confirmed in other experiments testing GMKL on pure sine wave data). Again, GMKL is able to ignore the source that is not useful to it. (It looks like GMKL on concatenated data is more effective when one source is a sine wave than when both are XORs as in Experiment 2, but this is because this experiment uses a 3D XOR rather than the 5D XOR of Experiment 2).

Conclusion

This paper investigated the effect of decision boundary complexity, unreliable sources, and uninformative dimensions on both a specialized data integration method as well as traditional classifiers applied to concatenated data. Generalized Multiple Kernel Learning (GMKL) was found to outperform all other classifiers on all experiments detailed here, because it was able to build separate representations of each source. This performance above and beyond the other classifiers was only present when GMKL was allowed to use separate kernels for each source, and not when it was forced to use concatenated data. This suggests that data integration techniques that take advantage of disparate sources in different ways have huge potential over concatenating data together. Treating each source separately allows for a reduced effective dimensionality, as well as robustness to unreliable sources and extra uninformative dimensions. In general, it is expected that data integration techniques will achieve better classification and better analysis than naive techniques that concatenate disparate datasets together.

Future Work

While these experiments show lots of promise, there are still many questions to answer. Data integration, data simulation, and benchmarking of machine learning techniques are all fields in need of further investigation. In particular, the behavior of machine learning and data integration techniques as functions of properties of the data and its decision boundaries are very poorly understood. Some fruitful further questions for research include:

- **Shared Information Across Sources:** An important component of real world multi-source datasets is that information is shared across sources. That is, the sources aren't completely independent of each

other, but both give insights into similar underlying hidden variables. The current data generation network model presented above does not have a way to systematically vary the information that is shared between both sources, but it would be very interesting to investigate if and how this property of the data affected the performance of classifiers.

- **Neural Networks:** Artificial neural networks are very powerful classifiers that are growing in popularity. They are an extremely extensible and flexible framework, with specialized versions existing for specific problems within the machine learning. They were not included in these experiments because they are too flexible. Neural networks have many more parameters than most popular classifiers, so it was not known how to pick an appropriate network architecture and topology for comparison against the other classifiers.

- **Theoretical Framework of Decision Boundaries:** We saw in Experiment 4 that GMKL performed well on the XOR dataset and not well on the Sine dataset. We would like to develop a theoretical framework that characterizes the aspects of decision boundaries that enable or disable a particular data integration technique from performing well on a dataset. We would like to come up with a concrete method for determining the best choice of data integration technique for a specific decision boundary.

- **More Computationally Expensive Testing:** All of the experiments in this paper ran with $N = 5000$ data points. For really thorough results, it'd be better to run these experiments for several values of $N \in \{1000, 5000, 1000, 2000, 50000, 10000\}$ as well as for more values of noise, more dimensions, more corruption levels, more width γ and regularization C of the SVMs, etc.

- **More varied simulated data:** This paper only used two types of decision boundaries, that of the sine wave and that of the N-dimensional XOR. There are many types of decision boundaries that are not represented here and many possible properties of data that aren't present here. Possible unimplemented ideas were to create separable data in $m < N$ dimensions and then to use (mostly) monotonic transformations to map the m dimensions into N-dimensional space, to create N dimensional separable data.

Appendix

Variable	Description
$M_{i,useful}$	Number of useful Data Dimensions for Data Source i
$M_{i,noisy}$	Number of Noisy Dimensions for Data Source i
N	Number of Data Points Generated (training and testing)
T	Ratio of Training to Testing
p_1, p_2	Probability of Data Source Corruption
σ_i	Variance of Gaussian Noise
C	Proportion of Data points in each class
k	Period of Sine Curve

Experiment 1: Data Dimension Scaling

XOR Dimensions (per source)	KNN	Random Forest	SVM	Concatenate + GMKL	GMKL
2	99.0	96.6	99.0	99.9	99.9
3	86.2	79.7	83.6	97.5	99.8
4	70.1	68.7	68.9	78.6	97.5
5	60.8	58.1	61.4	66.2	96.4
6	57.8	55.0	60.2	64.6	83.7
7	58.1	54.8	57.2	63.5	74.9

Experiment 2: Corrupted XOR Sources

p_1	p_2	KNN	Random Forest	SVM	Concatenate + GMKL	GMKL
0.0	0.0	61.5	58.6	62.6	67.7	94.7
0.0	0.1	61.2	56.8	62.0	65.8	92.3
0.0	0.2	57.6	58.3	59.5	60.4	88.8
0.0	0.3	55.2	54.9	57.0	58.6	88.6
0.0	0.4	57.3	56.7	59.7	60.8	91.2
0.0	0.5	57.3	57.4	57.4	56.0	91.1
0.1	0.1	58.2	55.7	59.5	62.5	85.5
0.1	0.2	54.7	53.2	57.3	61.3	80.7
0.1	0.3	54.3	53.9	55.1	56.9	79.2
0.1	0.4	54.8	53.5	54.7	55.0	81.6
0.1	0.5	54.9	54.3	53.8	60.4	79.2
0.2	0.2	55.3	50.8	55.4	61.0	68.3
0.2	0.3	54.2	51.6	55.6	55.4	71.1
0.2	0.4	50.9	51.5	49.7	56.5	63.1
0.2	0.5	50.4	50.7	52.5	54.9	59.6
0.3	0.3	52.4	48.1	50.9	54.3	63.6
0.3	0.4	52.2	51.7	51.2	54.0	62.4
0.3	0.5	53.3	50.0	48.6	52.3	56.6
0.4	0.4	51.0	52.4	53.1	50.0	53.7
0.4	0.5	50.7	48.4	50.8	48.4	52.2
0.5	0.5	49.1	49.7	49.0	48.8	48.5

Experiment 3: Noise Dimensions

XOR Dimensions (per source)	Noise Dimensions (per source)	KNN	Random Forest	SVM	Concatenate + GMKL	GMKL
3	0	86.7	80.0	84.3	97.5	99.8
3	1	81.7	77.3	76.3	79.7	100.0
3	2	76.0	78.2	74.3	74.5	98.9
3	3	73.6	77.0	74.8	74.2	95.3
3	6	61.7	69.6	69.4	74.4	77.0
5	0	61.0	60.4	63.5	66.2	96.4
5	1	58.4	56.6	61.4	66.7	72.3
5	3	57.5	54.2	59.5	67.2	70.0
5	5	57.0	54.2	58.4	66.3	65.6
5	10	54.7	54.3	58.3	64.0	65.1
7	0	57.7	54.9	60.2	63.5	74.9
7	1	52.4	52.1	53.8	64.2	66.2
7	4	55.3	53.8	56.5	59.8	62.1
7	7	54.8	54.1	56.5	59.7	59.9
7	14	51.6	52.1	54.5	61.9	59.4

References

[NBM14]   Elias Chaibub Neto, J Christopher Bare, and Adam A Margolin. Simulation studies as designed experiments: the comparison of penalized regression models in the large p, small n setting. *PloS one*, 9(10):e107957, 2014.

[YXY⁺11] H. Yang, Z. Xu, J. Ye, I. King, and M. R. Lyu. Efficient sparse generalized multiple kernel learning. *IEEE Transactions on Neural Networks*, 22(3):433–446, March 2011.